

# CS 295: Modern Systems

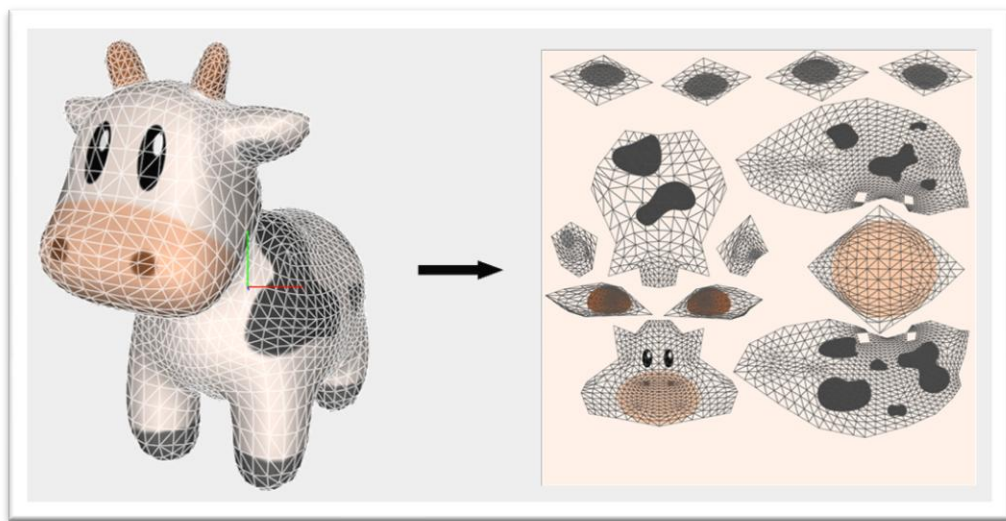
## GPU Computing Introduction




Sang-Woo Jun  
Spring 2019

# Graphic Processing – Some History

- ❑ 1990s: Real-time 3D rendering for video games were becoming common
  - Doom, Quake, Descent, ... (Nostalgia!)
- ❑ 3D graphics processing is immensely computation-intensive



Texture mapping



Shading

# Graphic Processing – Some History

- ❑ Before 3D accelerators (GPUs) were common
- ❑ CPUs had to do all graphics computation, while maintaining framerate!
  - Many tricks were played



## Doom (1993) : “Affine texture mapping”

- Linearly maps textures to screen location, disregarding depth
- Doom levels did not have slanted walls or ramps, to hide this

# Graphic Processing – Some History

- ❑ Before 3D accelerators (GPUs) were common
- ❑ CPUs had to do all graphics computation, while maintaining framerate!
  - Many tricks were played



Quake III arena (1999) : “Fast inverse square root”  
magic!


```
float Q_rsqrt( float number )
{
    const float x2 = number * 0.5F;
    const float threehalfs = 1.5F;

    union {
        float f;
        uint32_t i;
    } conv = {number}; // member 'f' set to value of 'number'.
    conv.i = 0x5f3759df - ( conv.i >> 1 );
    conv.f *= ( threehalfs - ( x2 * conv.f * conv.f ) );
    return conv.f;
}
```




# Introduction of 3D Accelerator Cards

- ❑ Much of 3D processing is short algorithms repeated on a lot of data
  - pixels, polygons, textures, ...
- ❑ Dedicated accelerators with simple, massively parallel computation



A Diamond Monster 3D, using the Voodoo chipset (1997)  
(Konstantin Lanzet, Wikipedia)







# Peak Performance vs. CPU

	Throughput	Power	Throughput/Power
Intel Skylake	128 SP GFLOPS/4 Cores	100+ Watts	~1 GFLOPS/Watt
NVIDIA V100	15 TFLOPS	200+ Watts	~75 GFLOPS/Watt

Also,



# System Architecture Snapshot With a GPU (2019)




Lots of moving parts!



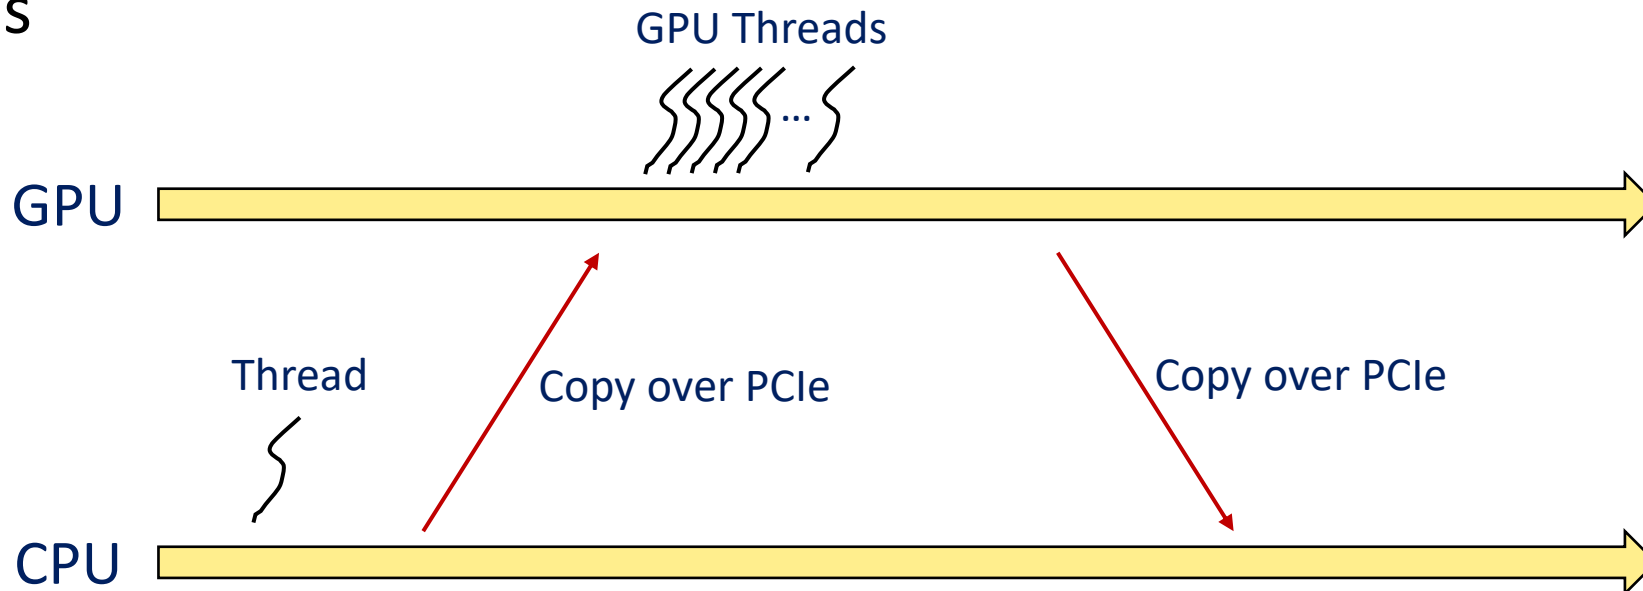
# High-Performance Graphics Memory

- ❑ Modern GPUs even employing 3D-stacked memory via silicon interposer
  - Very wide bus, very high bandwidth
  - e.g., HBM2 in Volta



# Massively Parallel Architecture For Massively Parallel Workloads!


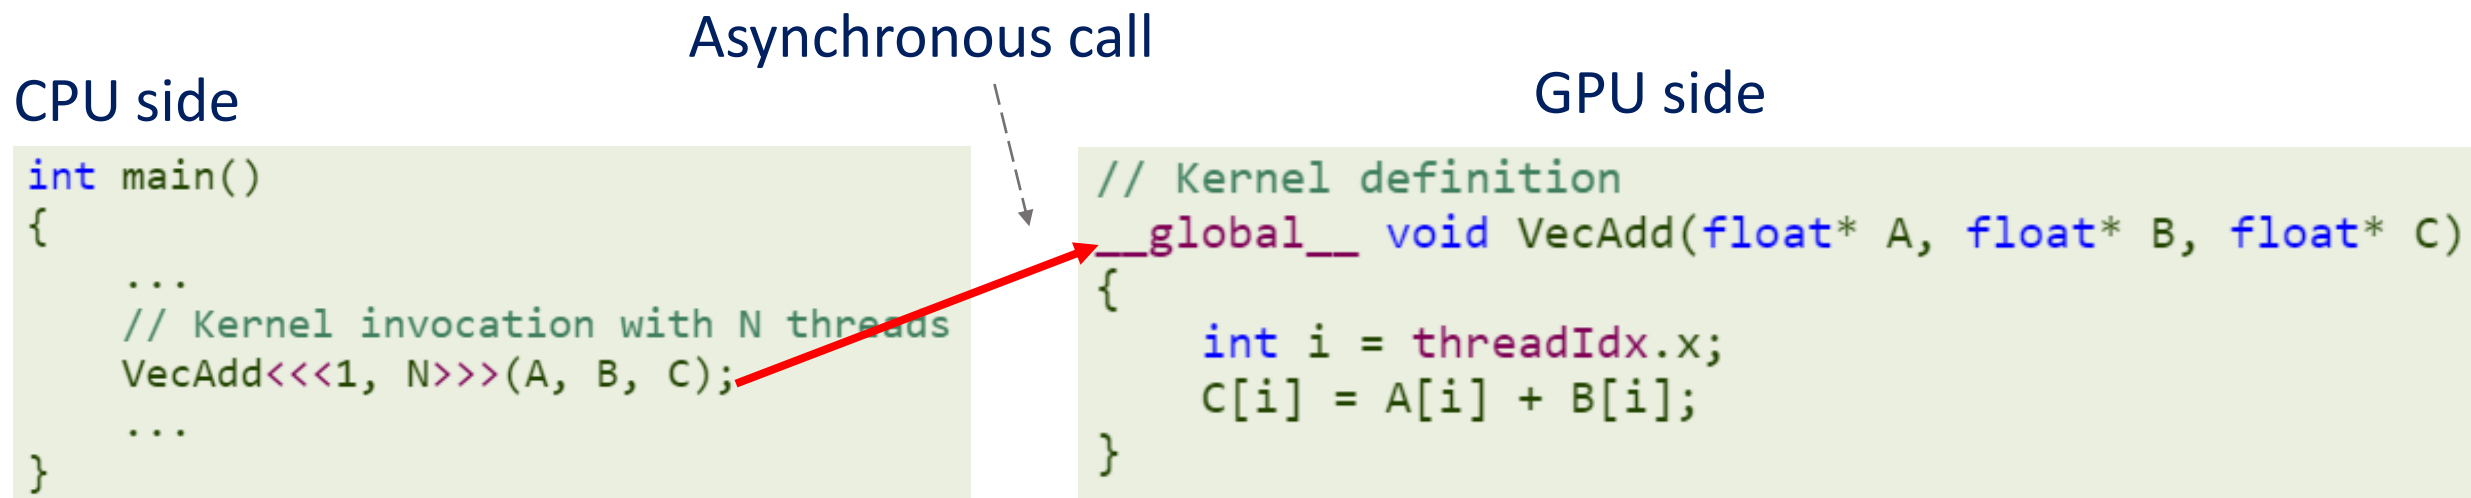
- ❑ NVIDIA CUDA (Compute Uniform Device Architecture) – 2007
  - A way to run custom programs on the massively parallel architecture!
- ❑ OpenCL specification released – 2008
- ❑ Both platforms expose synchronous execution of a massive number of threads



# CUDA Execution Abstraction

- ❑ Block: Multi-dimensional array of threads
  - 1D, 2D, or 3D
  - Threads in a block can synchronize among themselves
  - Threads in a block can access shared memory
  - CUDA (Thread, Block)  $\sim$  OpenCL (Work item, Work group)
- ❑ Grid: Multi-dimensional array of blocks
  - 1D or 2D
  - Blocks in a grid can run in parallel, or sequentially
- ❑ Kernel execution issued in grid units
- ❑ Limited recursion (depth limit of 24 as of now)

# Simple CUDA Example





# Simple CUDA Example

```
int main()
{
    ...
    // Kernel invocation with N threads
    VecAdd<<<1, N>>>(A, B, C);
}
```

1 block

N threads per block

Should wait for kernel to finish

\_\_global\_\_:  
In GPU, called from host/GPU  
\_\_device\_\_:  
In GPU, called from GPU  
\_\_host\_\_:  
In host, called from host

N instances of VecAdd spawned in GPU

```
// Kernel definition
__global__ void VecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x;
    C[i] = A[i] + B[i];
}
```

Only void allowed


Which of N threads am I?  
See also: blockIdx

One function can  
be both

# More Complex Example: Picture Blurring

- ❑ Slides from NVIDIA/UIUC Accelerated Computing Teaching Kit
- ❑ Another end-to-end example  
<https://devblogs.nvidia.com/even-easier-introduction-cuda/>
- ❑ Great! Now we know how to use GPUs – Bye?

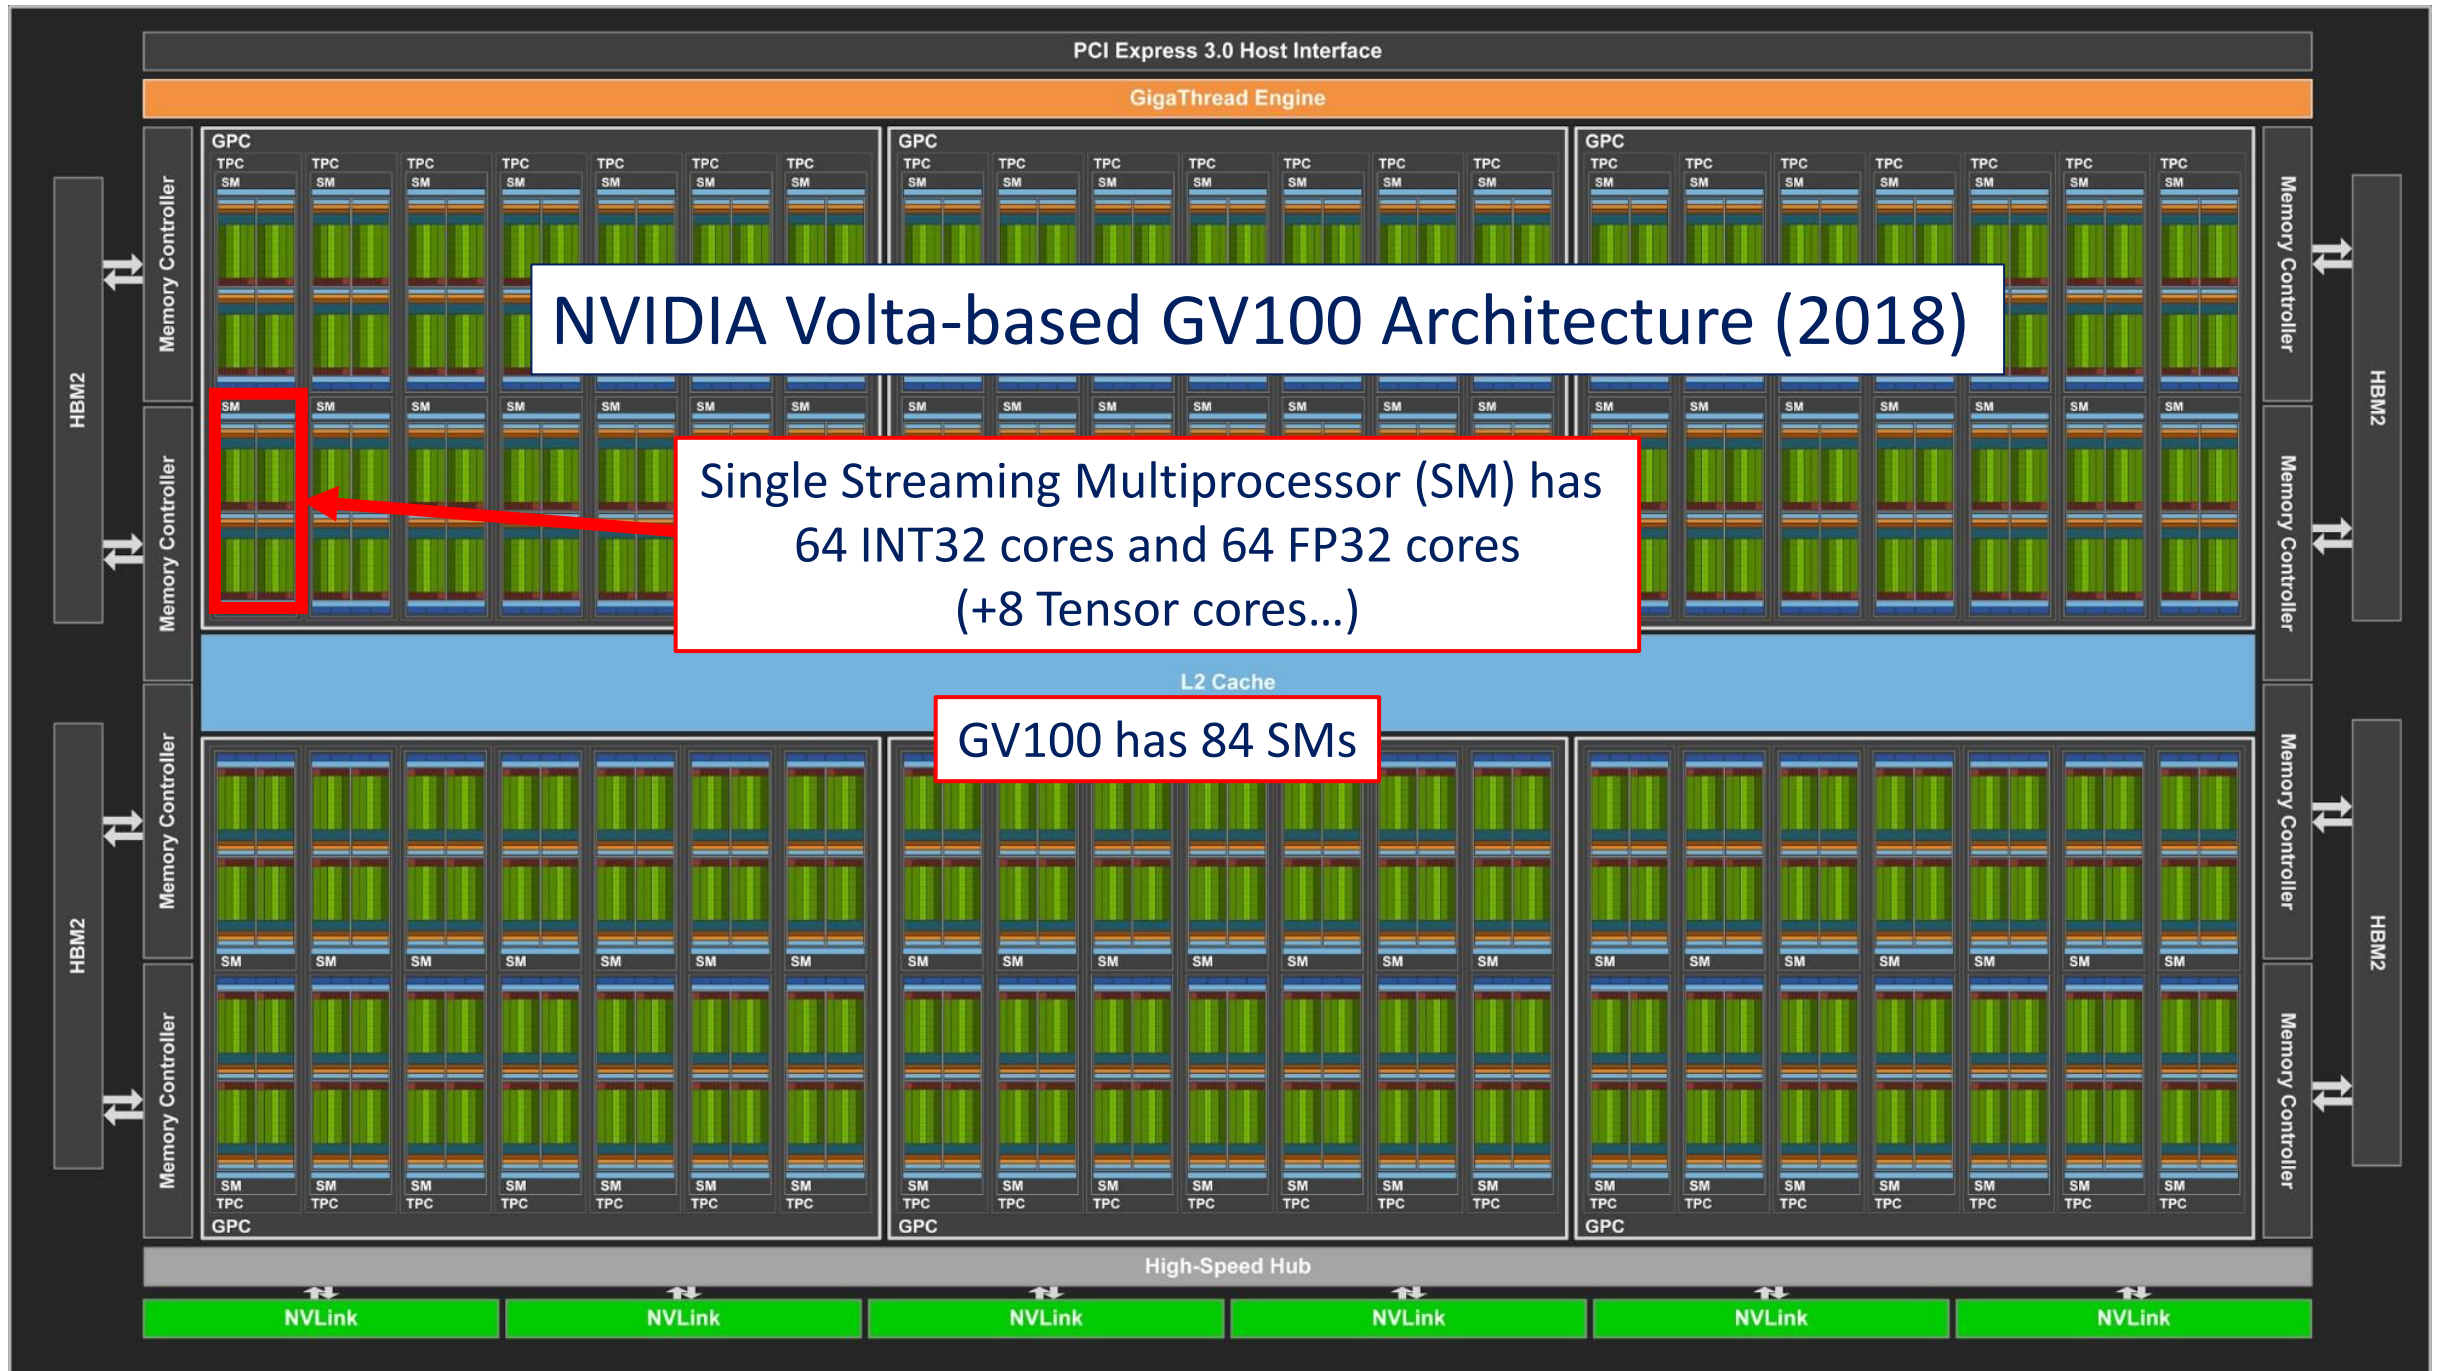
# Matrix Multiplication Performance Engineering



No faster than CPU

Results from NVIDIA P100

Architecture knowledge is needed (again)





# Volta Execution Architecture

- ❑ 64 INT32 Cores, 64 FP32 Cores, 4 Tensor Cores, Ray-tracing cores..
  - Specialization to make use of chip space...?
- ❑ Not much on-chip memory per thread
  - 96 KB Shared memory
  - 1024 Registers per FP32 core
- ❑ Hard limit on compute management
  - 32 blocks AND 2048 threads AND 1024 threads/block
  - e.g., 2 blocks with 1024 threads, or 4 blocks with 512 threads
  - Enough registers/shared memory for all threads must be available (all context is resident during execution)

More threads than cores – Threads interleaved to hide memory latency



# Resource Balancing Details

- ❑ How many threads in a block?
- ❑ Too small: 4x4 window == 16 threads
  - 128 blocks to fill 2048 thread/SM
  - SM only supports 32 blocks -> only 512 threads used
    - SM has only 64 cores... does it matter? Sometimes!
- ❑ Too large: 32x48 window == 1536 threads
  - Threads do not fit in a block!
- ❑ Too large: 1024 threads using more than 64 registers
- ❑ Limitations vary across platforms (Fermi, Pascal, Volta, ...)

# Warp Scheduling Unit

- ❑ Threads in a block are executed in 32-thread “warp” unit
  - Not part of language specs, just architecture specifics
  - A warp is SIMD – Same PC, same instructions executed on every core
- ❑ What happens when there is a conditional statement?
  - Prefix operations, or control divergence
  - More on this later!
- ❑ Warps have been 32-threads so far, but may change in the future

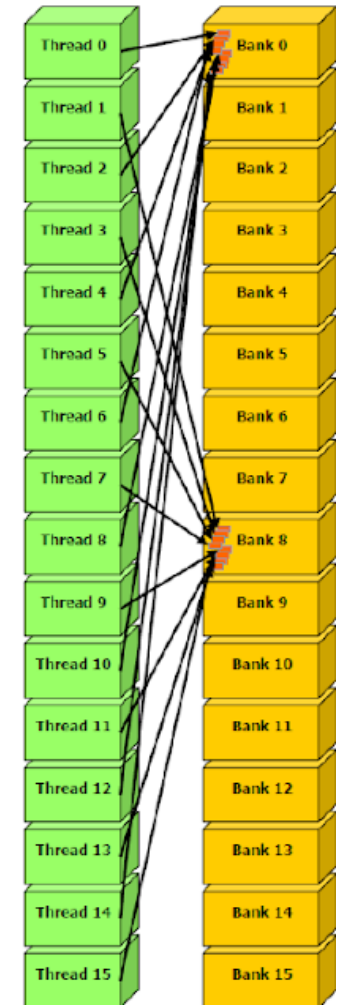
# Memory Architecture Caveats

## ❑ Shared memory peculiarities

- Small amount (e.g., 96 KB/SM for Volta) shared across all threads
- Organized into banks to distribute access
- Bank conflicts can drastically lower performance

## ❑ Relatively slow global memory

- Blocking, caching becomes important (again)
- If not for performance, for power consumption...



8-way bank conflict  
1/8 memory bandwidth